

And the docker compose will be very over welling , we need something more powerful that automates deployment , scaling of the containers .

Topics to be covered:

1. Introduction to k8s
2. Main k8s components -> pods , node, sVc , secret, deploy , Sts, ing. etc.
3. Minikube & kubectl - Local setup
4. Main kubectl commands -> interacting with k8s cluster.
5. YAML configuration file -> configure k8s components
6. Organising components using **namespaces**
7. Configuration connectivity **services** .
8. Make app available outside the cluster **Ingress**
9. Persistence of data in k8s using **volumes**
10. ConfigMap and secrets as volumes type
11. Deploy stateful app like databases using **StatefulSet**
12. Package manager tool for k8s - **helm**
13. Extending the k8s apis - **operator**

Also learn how to deploy microservices applications in k8s

Production and security best practices in k8s.

Authorisation in k8s- role based access control (RBAC)

## Lets start :)

K8S is a open source container orchestration framework or tool .

Developed by google

Helps u manage application which are made of 100's or 1000's of containers.

It helps u manage them in different container environment eg: physical , virtual , hybrid or cloud environment .

## What problem does k8s solves ?

Tread from monolithic to microservices , increased use of containers .

## What features does container orchestration offers ?

1. High availability or no downtime.
2. Scalability or high performance => which high load in the serves scale up vica versa!
3. Disaster recovery - back up and restore.

## Main k8s components :

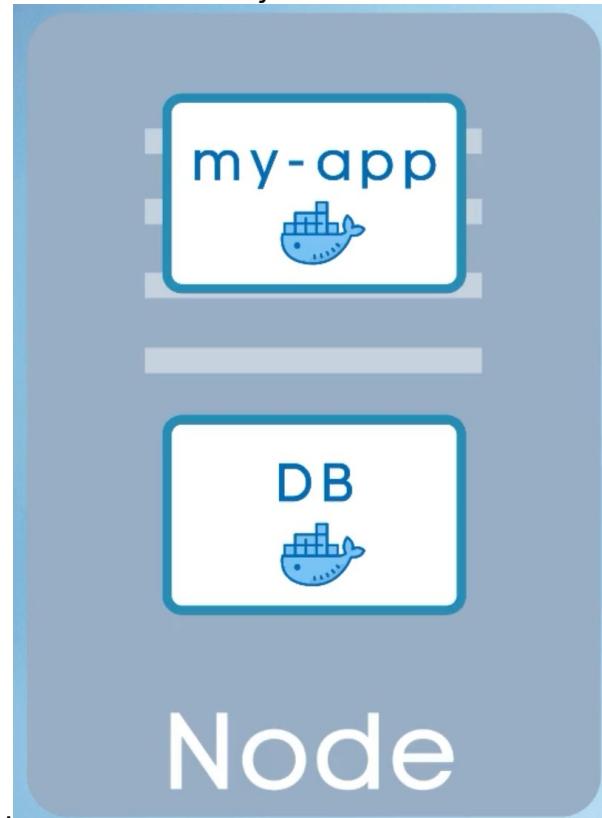
Pod , service , configMap, Deployment , Ingress , Secret , StatefulSet .

### 1. Node and pod !

**Node** = virtual or physical environment .

Pod is a smallest unit in kubernetes .

Abstraction over the container so that u don't need to interact with the container technology and only k8s will interact with its layer .



Usually 1 application per pod .

Here u have one server and two container pods running on it with a abstraction layer on top on them .

How these pods contact with each other -> each pods get a ip address.  
And each pod can communicate with each other using that internal ip address its not the public ip address.

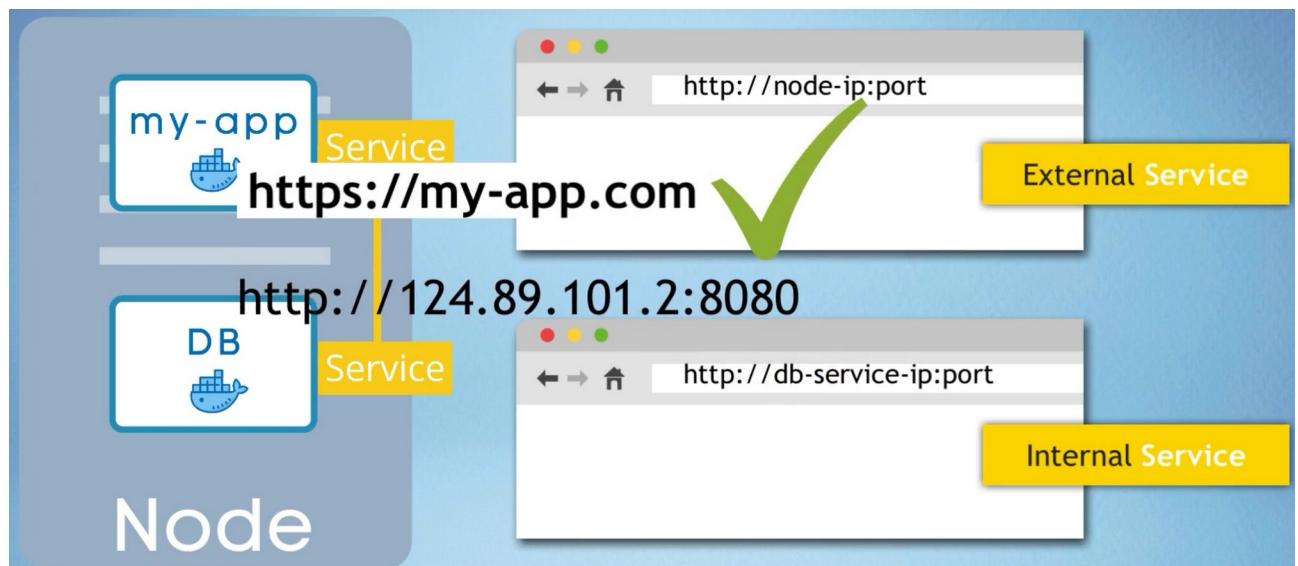
### **Pods are ephemeral!**

Means they can die easily because of some error like app crashed inside or the server node were pods are run out of the resources and a new one will be created on its place with a new IP address assigned to it .

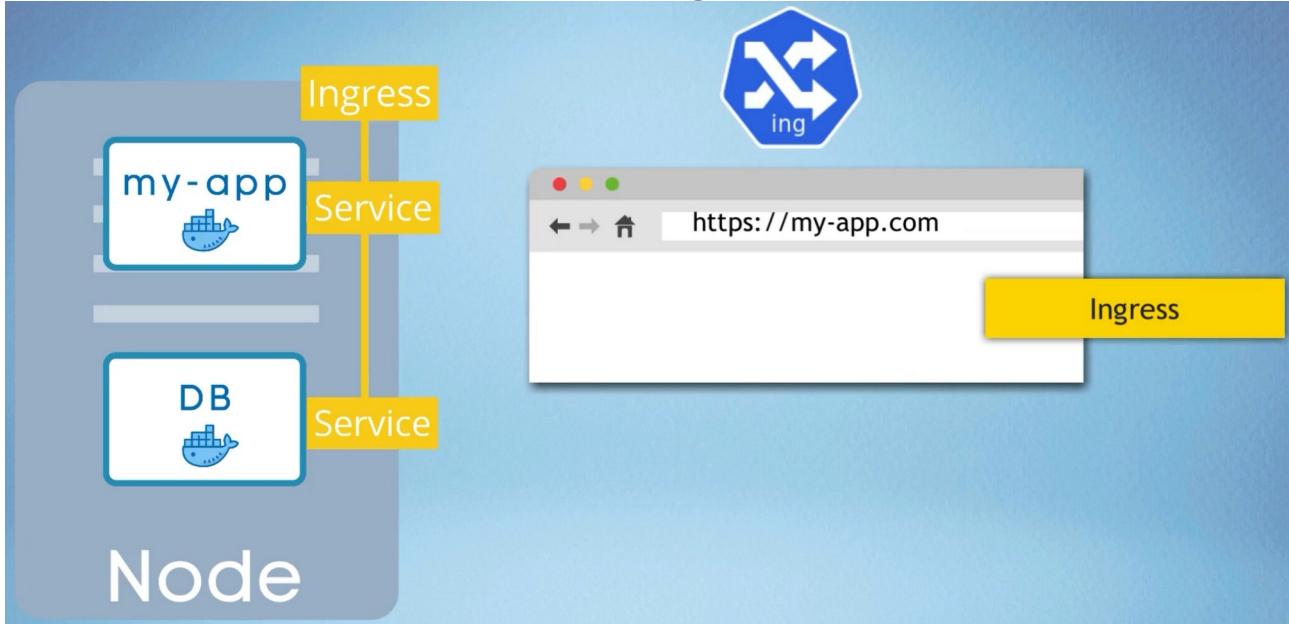
Now its very inconvenience to replace the ip address so because of that another component of k8s service is used .

Service : it is a permanent ip address , life cycle of service and pods are not connected .

Now u want your application to be accessible in the browser so for that u need a external service .



Instead of service the request goes through **Ingress** and it forward it to the service .



### **configMap and secret :**

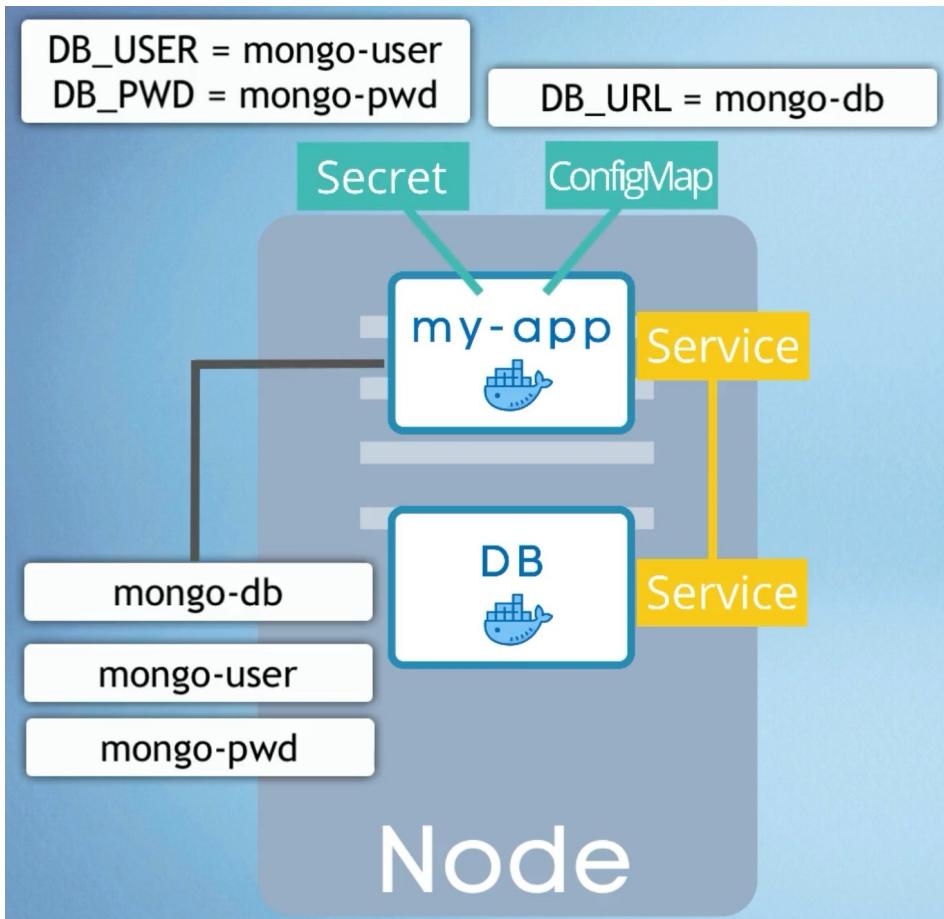
My app will talk to eg: mongo-db service but where do u configure all these -> it is usually database URL is in the build image of the application . So for example the endpoint of the service will change to oracle-db etc. so u would have to adjust that url in that application so u will have to rebuild the application with new version -> and u will have to push it to the repo -> pull that new image in your pod and restart every thing so it is little bit tedious to just change the URL of the database .

So for that. ConfigMap helps -> will contain external configuration of your application like URL of database and in k8s u just connect it to the pod .and that pod will get the data that the configMap contain , now u just change the information in the configMap and u don't need to build another image again.

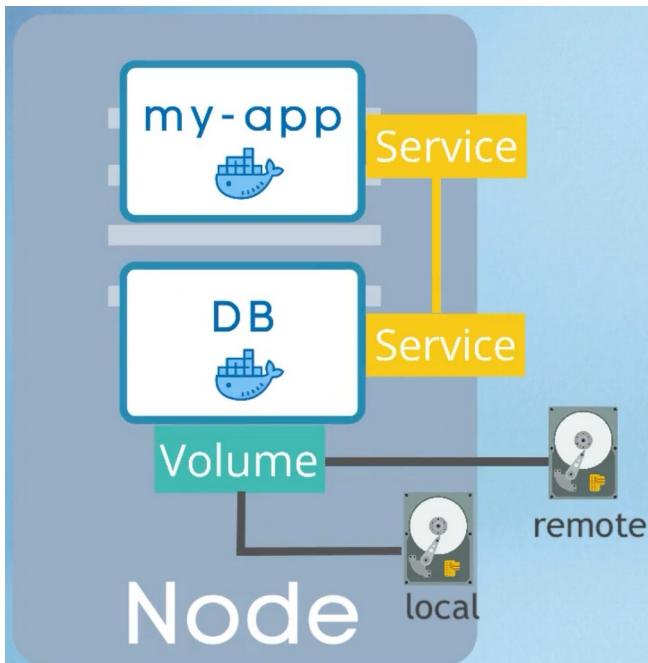
configMap is for the non - confidential data only.

So for the database user name and password u have another component in k8s **secret** .

We use 3rd party tool and deploy it to k8s those tools will encrypt the secret data k8s don't provide encrypting itself it used third party tools .



U can use configMap or secret inside your application pod using environment variable or as a properties file.



#### Volumes :

If the data base get restarted the data will be gone but using volumes u can persist the data => it attaches the physical storage in the hard drive to your pod it can be either in the local machine the same server node or it can be on the remote storage out side the k8's cluster .

K8s explicitly does not provide persistence of data .

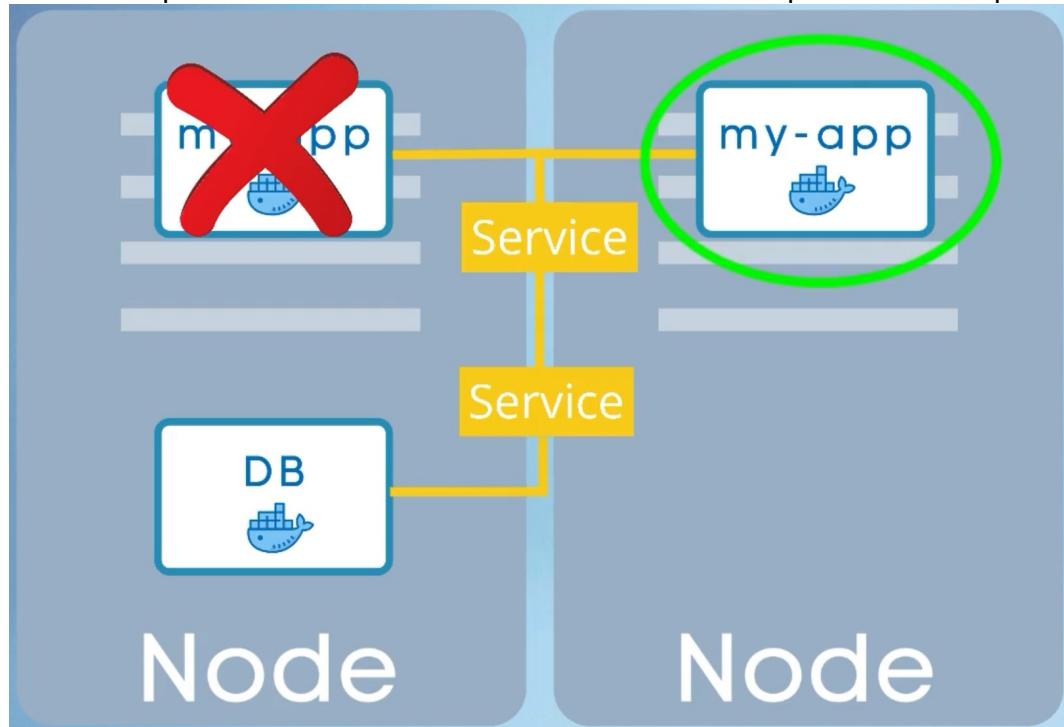
#### Deployment and statefulSets :

Till now every thing is good user can access our application through the browser. What happens if my application pod dies and user will have a downtime and can not reach the application which is very bad thing if it happen in production . So this is the advantage of distribution system and container so instead of relying on one applicant node we are replicating everything in on multiple server nodes

Service is like a permanent ip address with a dns name and service is also a load balancer .

U will not actually create a second pod rather than u will define a blueprint defining how many replicas u want to have . And that component is called **Deployment** . In practice u will not be creating pod u will be creating deployments because there u can scale up or down no of replicates of pods u need . **Pods were the layer of abstraction over the containers and deployment is another abstraction on top of pods** which makes it more convince to interact with pods and do some other configuration on them .

If the first pod will die the service will forward the request to other pod .



But what about replication Database -> we done replication DB using deployment . Reason of the is DB has a state meaning if we have a clone of DB they will have to access the same shared data storage and there u will need a mechanism recording which DB pod is currently writing or reading to that DB in order to avoid data inconsistency .

And that mechanism is offered by k8s component -> **statefulSet**

Eg: MYSQL , ELASTIC SEARCH etc all the stateful application should be creating using statefulset and not deployment .

Stateful set is just like deployments used to scale up or down database or replicate them and it makes sure no database inconsistency are offered .

Deploying DB application using statefulset in k8s cluster can be tedious so it more difficult so better practice is to host database outside the k8s cluster and just have a stateless application inside of the k8s cluster .

#### **DaemonSet:**

We want to collect logs from our cluster from each pod .

So we need a application which will run on each node and collect all the logs that pods are generating there

When we add node we need to adjust the replica no. in the deployment each time and when we delete a node we again need to adjust the replica no.

When deployment we cant also ensure all the pods are distributed equally among all the nodes.

So for that we have another k8s component which manages all these thing for u .

Daemon set automatically calculates how many replica u need based on the nodes and pods u have in the k8s cluster .

When add the node it will add the pods when delete node it will automatically remove pods (garbage collected)

Now u will not have to define the replicate count the deployment .

Automatically scale up and scale down .

Also guarantee one replica per pod .

#### **Wrap up the main components of k8s we have covered :**

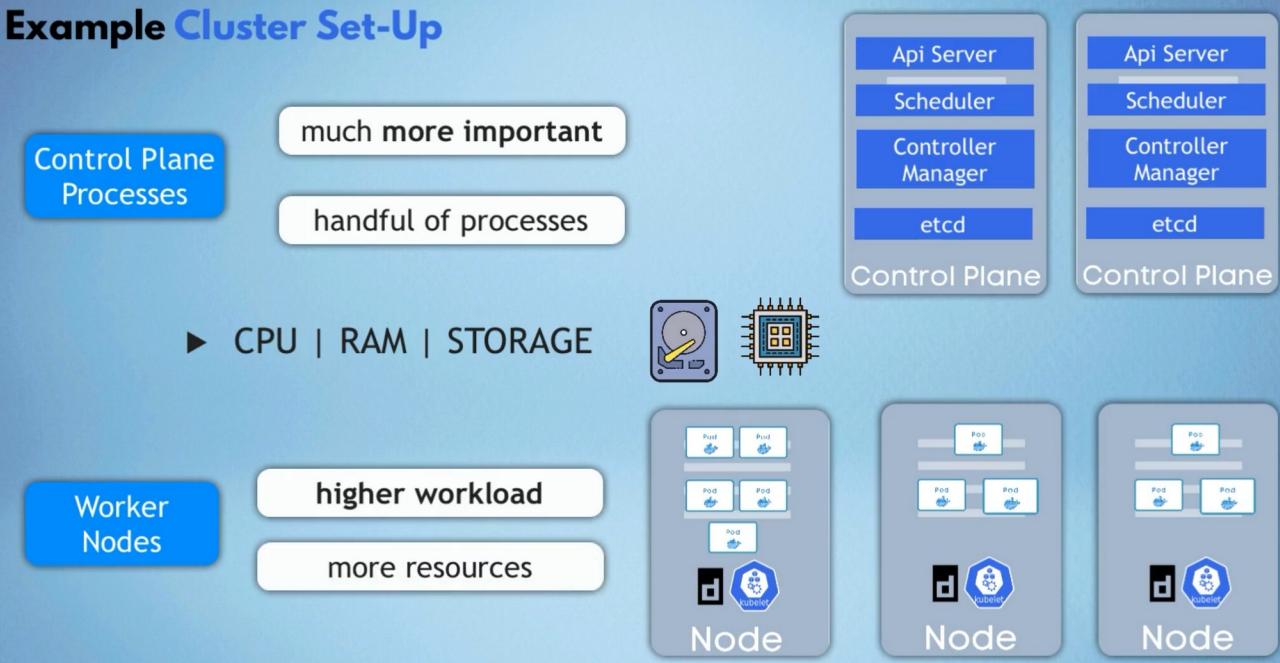
1. Pod : abstraction of container
2. Service : communication
3. Ingress: route traffic into cluster .
4. ConfigMap: external configuration
5. Secret: secret external configuration
6. Volumes : data persistence
7. Deployment : pod blueprint replication
8. Statefulset : DB pod blueprint replication
9. demonSet: automatic replica collect logs .

#### **Kubernetes Architecture :**

2-Types of nodes or machines

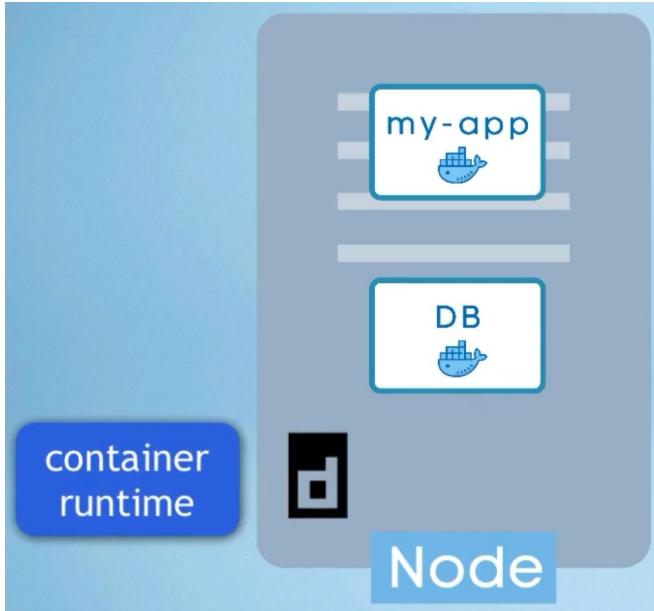
1. Control plane

## Example Cluster Set-Up



### 2. Worker node

What is the difference roles it have inside the k8s cluster how k8s does what it does how the cluster is self managed and self healing and automated how u as a operator on the k8s cluster has the must lesser work to done to mange all this .



### Node processes :

Nodes are the cluster servers that actually do the work there for some time called the worker node

Each node have multiple pods in it .

3 processes must be installed in each node .

**1process** is the contender run time eg: docker container , container D , crio-0 the most popular contender run time used today in most of the k8s cluster is container -D which is much more light weight than docker .

U can run docker images in any container run time there for u will see the docker icon in pods .

Container runtime need to be installed in every node .

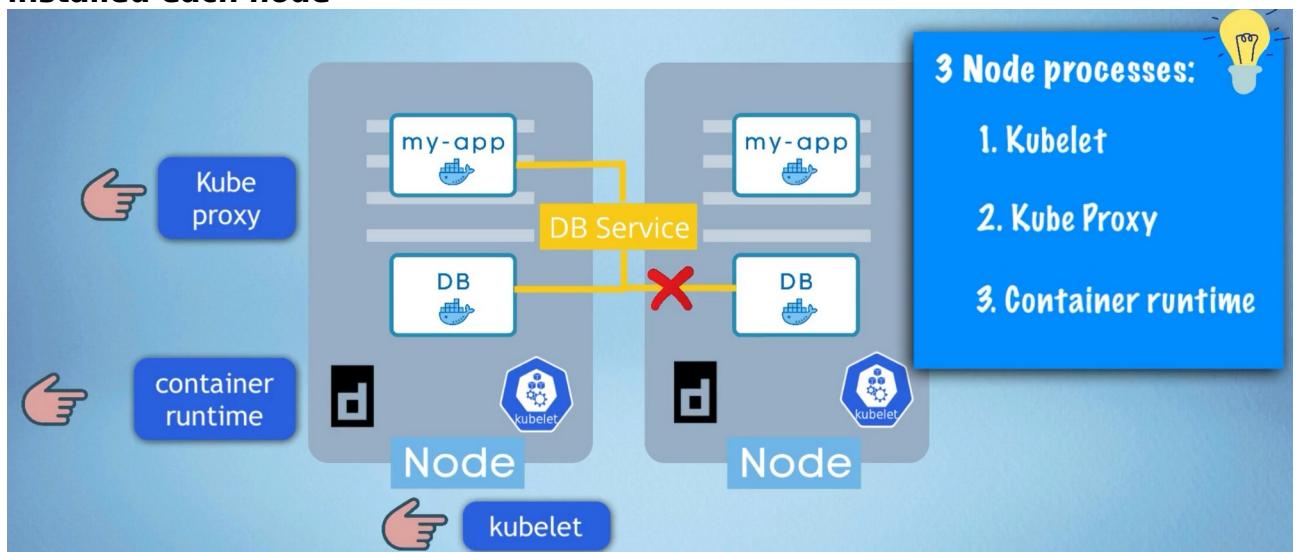
**2-process** Kubelet is the process of k8s which interact with both container and node . Kubelet will take the configuration starts the pod with a. container inside and assigning resources from that node to that container .

Usually k8s clusters are made up of multiple nodes which also must have container run time and kubelet services inside each node and u can have 100s of those node which would run other pods , container and replicates of existing pods **these nodes communicate with each other using services.**

Which is a load balancer which catches the request directed to the pod of application . And forward it to the respective pods .

**3- process Kube proxy** responsible for forwarding the requests from the services to pods .

**So container run time , kubelet and kubeproxy processes need to be installed each node**



**So , how do u interact with these cluster ?**

**How to know scheduled pod?**

**Monitor?**

**Reschedule/ restart ?**

**Join new node? Who do it join the cluster and join pods to it .**

**Ans is \_—> all of this process are managed by control plane nodes .**

**Control plane processes :**

4 - process run on every control plan node that control the cluster state and worker node as well.

**1process : api server** -> when as a user u want to deploy new application in k8s cluster u interact with Api server through UI or CLI . Api server is a like a cluster gateway also act as a gatekeeper for authentication .

So when every u want to schedule new pods deploy new applications >create new services or new other components u have to contact to the api server on the control plan node and api server than validates your request and if every thing is fine then it

will forward your request to other processes . In order to schedule the pod or create new component that you requested .

Also if you want to query the states of your cluster health etc . You make request to the api server and it gives back you the response .

Which is good for security as you have only one entry point to the cluster .

**2process : scheduler :->** scheduler will schedule which pod to deploy on which worker node or next component will be scheduled to which worker node.

**Process that actually does the scheduling that starts that pod with a container is the kubelet , so kubelet get the request from the scheduler and start a pod on the node .**

**3process : controller manager ->** when pods die on a node then we have to reschedule those pods again . Controller manager will detect cluster state changes . And try to recover the cluster state as soon as possible .

Controller manager request scheduler and scheduler decides based on the resource left calculation which worker node to schedule the pod .

**4process: etcd ->** etcd (key-value storage) is a cluster brain .

Which record every change in the cluster eg: new pod scheduled when a pod dies all of these changes get updated in the key value store .

All of the other process controller manager and scheduler works because of etcd data .

How does the scheduler do know which resources are available on each worker node .?

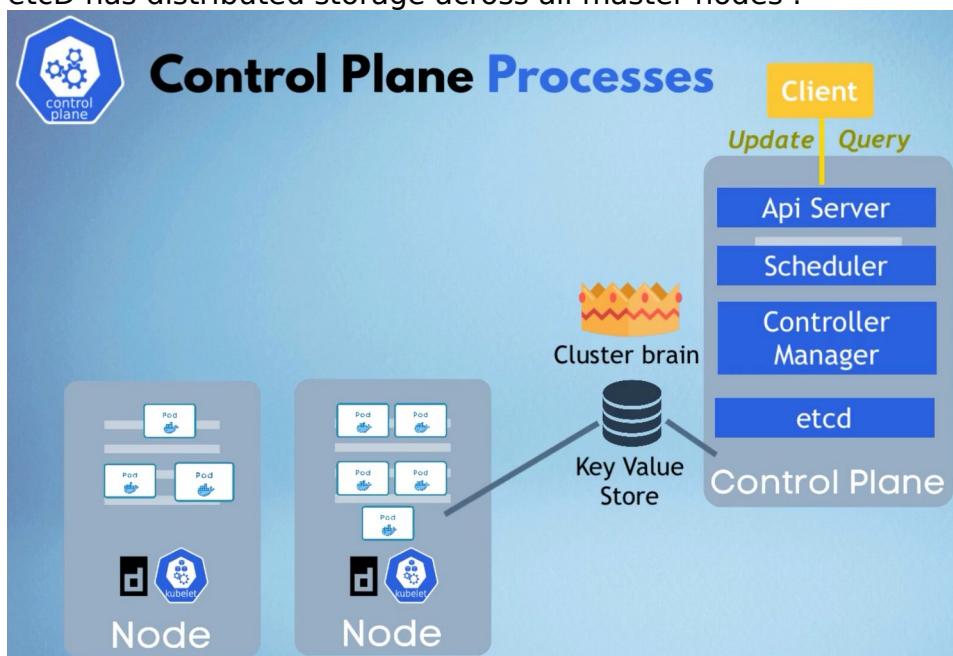
How does controller manager know which cluster state changes ?  
Is cluster health ok query request to the api ?

All of these information is stored in the etcd .

What is not stored in the actual etcd is the actual application data .

In practice k8s cluster is made up of multiple control planes , api server is load balanced .

etcd has distributed storage across all master nodes .



### **Cluster setup :**

In a very small cluster u would probably have 2 control plane node and 3 worker node .  
Console plane process are more important but they have lower work load and have handful of processes to run there they required lower resources like cpu , ram , storage :

On other side worked nodes has higher workload therefore required more resources .

As the complexity increased u can add more control plane and worker nodes on your cluster to make the application requirement full fill .

### **Minikube and kubeCtl :**

#### **What is Minikube ?**

So in a production cluster generally u will have multiple control plane node and worker nodes .

Separate virtual and physical machines which each represent a node .

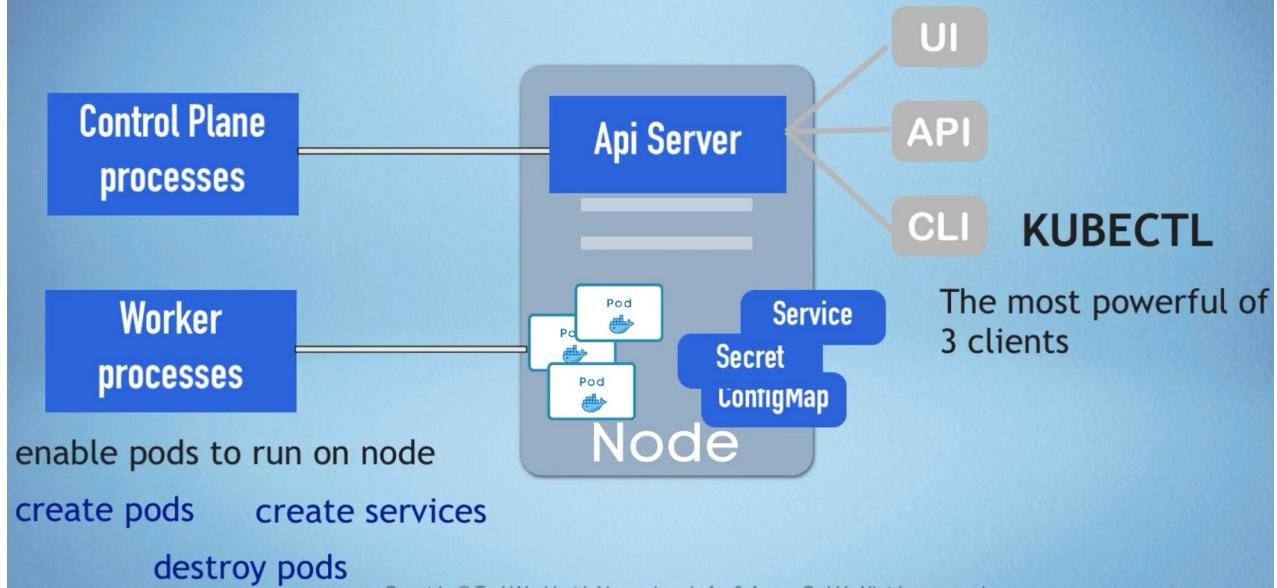
So if u will have to setup all this manually it should be very difficult as u will lack resources cpu , memory , ram etc .

So for that u have a open source tools called Minikube ->  
Minikube is a one node cluster where the control plane and worker process run on the one node.

And that node will have docker control run time preinstalled .

## What is kubeCTL ?

### What is kubectl?



So now u have this virtual node on your local machine that represent Minikube u need a way to interact with that cluster like creating pods and other k8s components on Minikube and u can do it via kubeCTL which is a command line tool for k8s .

Minikube run both control plan and worker processes .

One control plan process ie : API server is the main entry point for k8s cluster (enables interaction with cluster ) with different clients like UI , api , command line tool which is kubeCTL and kubeCTL is most powerful of all the clients .

Then worked nodes will enable pod to run on node -> create pods -> destroy pods or create services etc .

kubeCTL can interact with any type of k8s cluster eg: Minikube cluster , cloud cluster etc.

## Installation and creation of Minikube cluster :

Visit official documentation of Minikube .

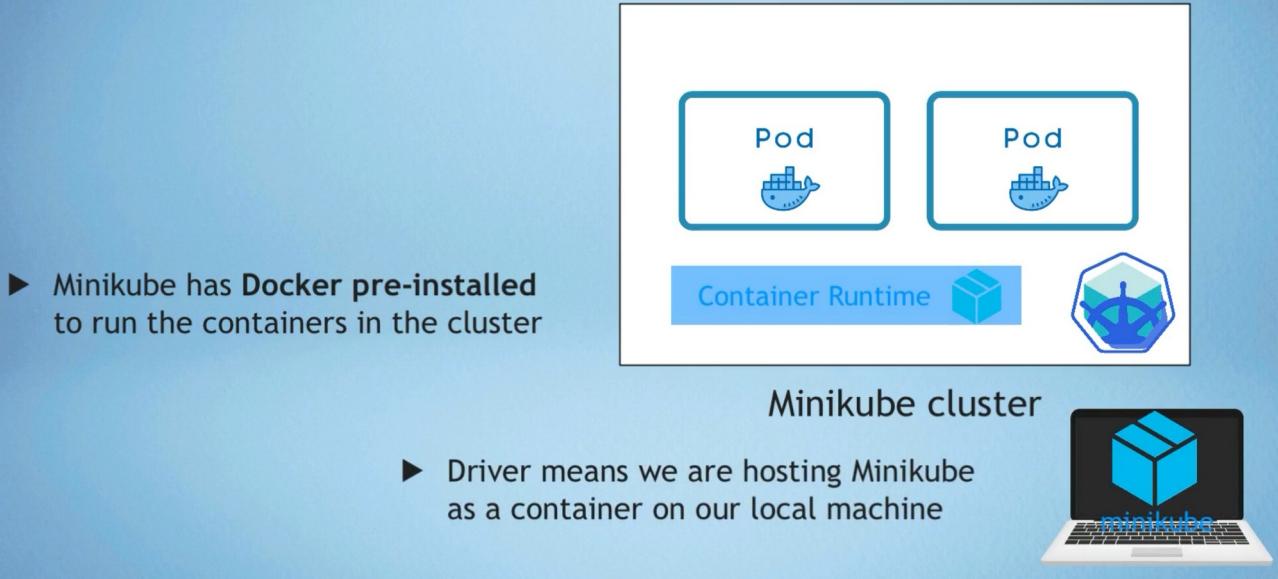
U can run Minikube either a container or a virtual machine on your laptop.  
So we need a virtual machine or container tool installed on our laptop eg: docker , container -D , cri\_O , vm ware , oracle virtual box etc .

For Mac OS - m1 - air ->

1. brew install minikube
2. Start your cluster -> minikube start —driver docker
3. Open the driver page to see the list of supported container or virtual machine tools .

4. Docker is a preferred driver to run Minikube on all OS .

## 2 Layers of Docker



**We have two layers 1. Minikube running as a docker container and docker inside Minikube to run our application container .**

5. When docker desktop installed open it .
6. To check the status of your cluster -> Minikube status .
7. Minikube has kubeCTL as dependency -> now u can interact with Minikube using kubeCTL
8. Kubectl get nodes -> get status of node .

## Kubectl CLI ...for configuring the Minikube cluster

```
 1. kubectl get nodes
 2. kubectl get pod
 3. kubectl get services
 4. kubectl create -h // to create components -h is help
 5. U will not see pods in above commands because pods are the smallest units in k8s cluster .
```

## Minikube CLI ...for start up/deleting the cluster

### Main kubectl command :

Create and debug pods in Minikube cluster :

1. Kubectl get nodes
2. Kubectl get pod
3. Kubectl get services
4. Kubectl create -h // to create components -h is help
5. U will not see pods in above commands because pods are the smallest units in k8s cluster .

Usually in practice u are not working with pods directly . There is a abstraction layer over the pods ie : deployment .

So this is what we will create and it will create pod under neath .

### Creating a nginx deployment ->

: kubectl create deployment nginx-depl --image=nginx // it will download the latest image of nginx from docker hub .

```
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl create deployment nginx-depl --image=nginx
deployment.apps/nginx-depl created
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get deployment
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
nginx-depl  1/1     1           1           59s
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
nginx-depl-85c9d7c5f4-c87zm  1/1     Running   0          2m32s
```

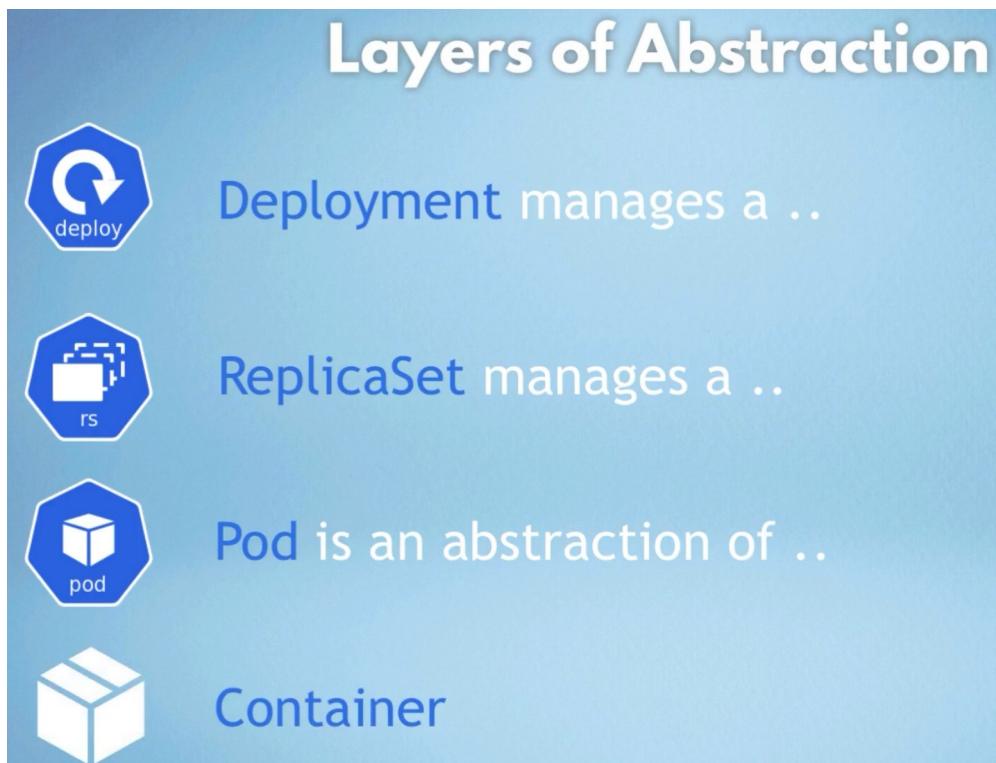
**Between deployment and pod there is another layer called replicaset**

: kubectl get replicaset

```
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get replicaset
NAME      DESIRED   CURRENT   READY   AGE
nginx-depl-85c9d7c5f4  1         1         1       4m21s
```

Deployment manages a -> replicaset manages a -> pod is a abstraction layer to container \_>container .

Everthing below deployment is managed by k8s .



```
: kubectl edit deployment nginx-depl // autogenerated file with default values.  
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl edit deployment nginx-depl  
Edit cancelled, no changes made.
```

Deployment Service

1) metadata

```
! nginx-deployment.yaml ×  
1 apiVersion: apps/v1  
2 kind: Deployment  
3 metadata:  
4 | name: nginx-deployment  
5 > labels: ...  
7 spec:  
8 | replicas: 2  
9 > selector: ...  
12 > template: ...
```

```
! nginx-service.yaml ×  
1 apiVersion: v1  
2 kind: Service  
3 metadata:  
4 | name: nginx-service  
5 spec:  
6 > selector: ...  
8 > ports: ...  
12
```

### Debugging pods :

Kubectl logs <podName>

```
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl create deployment mongo-deployment --image=mongo  
deployment.apps/mongo-deployment created  
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get pods  
NAME READY STATUS RESTARTS AGE  
mongo-deployment-5dfbb89958-26dvq 0/1 ContainerCreating 0 23s  
nginx-depl-85c9d7c5f4-c87zm 1/1 Running 0 16m  
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl describe pod mongo-deployment-5dfbb89958-26dvq  
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get pod  
NAME READY STATUS RESTARTS AGE  
mongo-deployment-5dfbb89958-26dvq 1/1 Running 0 2m20s  
nginx-depl-85c9d7c5f4-c87zm 1/1 Running 0 17m
```

after some time container inside the pod will run automatically .

Terminal of mongo DB application container ->

Kubectl exec -it(interactive terminal ) pod\_name — bin/bash

```
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl exec -it mongo-deployment-5dfbb89958-26dvq -- bin/bash
root@mongo-deployment-5dfbb89958-26dvq:/# ls
bin  data  docker-entrypoint-initdb.d  home    lib   mnt  proc  run  srv  tmp  var
boot dev   etc           js-yaml.js  media  opt  root  sbin  sys  usr
root@mongo-deployment-5dfbb89958-26dvq:/# exit
exit
```

**Delete deployment ,apply configuration file . :**

```
((base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get deployment
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
mongo-deployment  1/1      1          1         9m2s
nginx-depl     1/1      1          1        24m
((base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get pod
NAME          READY  STATUS    RESTARTS  AGE
mongo-deployment-5dfbb89958-26dvq  1/1    Running   0         9m16s
nginx-depl-85c9d7c5f4-c87zm       1/1    Running   0         24m
((base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl delete deployment mongo-deployment
deployment.apps "mongo-deployment" deleted
((base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get pod
NAME          READY  STATUS    RESTARTS  AGE
nginx-depl-85c9d7c5f4-c87zm       1/1    Running   0         25m
((base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get replicaset
NAME          DESIRED  CURRENT  READY  AGE
nginx-depl-85c9d7c5f4            1        1      1    26m
((base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl delete deployment nginx-depl
deployment.apps "nginx-depl" deleted
(base) vishaldwivedi@vishals-MacBook-Air-2 ~ % kubectl get replicaset
No resources found in default namespace.
```

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: nginx-deployment
5  >  labels: ...
7  spec:
8    replicas: 2
9  >  selector: ...
12  template:
13    metadata:
14      labels:
15        app: nginx
16    spec:
17      containers:
18        - name: nginx
19          image: nginx:1.16
20          ports:
21            - containerPort: 8080
```

- Since it's an abstraction over Pod, we have the **Pod configuration inside Deployment configuration**
- Own "metadata" and "spec" section
- Blueprint for Pod

All the crud operation happens at the deployment level every thing under neath just followed automatically .

**U will usually work with k8s configuration files as -> there can be many option in the command line that u can put which is tedious .**

```
% kubectl create deployment name image option1 option2
```

Kubectl apply will take the configuration file as a parameter and does what every u have written there .

**Creating configuration file .=>**

Eg: Kubectl apply -f nginx-deployment.yaml

So first creating this file :

: touch nginx-deployment.yaml

### Labels & Selectors

**Labels**

- Labels are **key/value pairs** that are attached to resources, such as Pods.
- Used to specify identifying attributes that are meaningful and relevant to users

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: nginx-deployment
5    labels:
6      app: nginx
7  spec:
8    replicas: 2
9    selector:
10   matchLabels:
11     app: nginx
12   template:
13     metadata:
14       labels:
15         app: nginx

```

**Label Selectors**

- Labels do not provide uniqueness
- Via **selector** the user can identify a set of resources

**Connecting Services to Deployments**

```

1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: nginx-service
5  spec:
6    selector:
7      app: nginx
8    ports: ...

```

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: nginx-deployment
5    labels:
6      app: nginx
7  spec:
8    replicas: 2
9    selector:
10   matchLabels:
11     app: nginx
12   template:
13     metadata:
14       labels:
15         app: nginx

```

**All the commands used till now ->**

**install hyperkit and minikube**

brew update

brew install

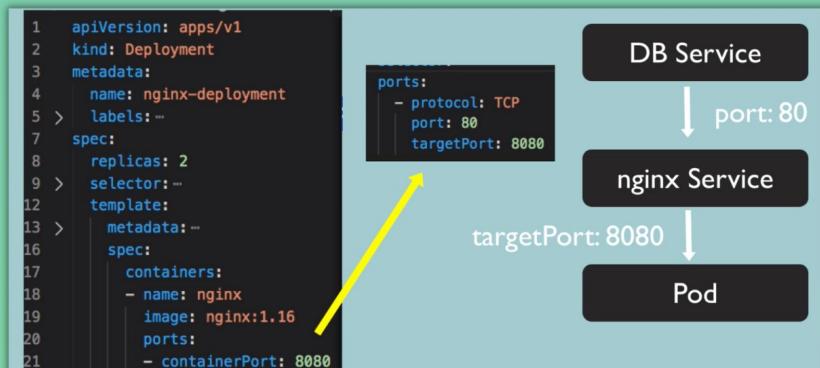
## Ports in Service and Pod

- In Service component you need to specify:

- port** = the port where the service itself is accessible
- targetPort** = port, the container accepts traffic on

- In Deployment component:

- containerPort** = port, the container accepts traffic on



hyperkit

brew install minikube

kubectl

The screenshot shows a code editor with two tabs open. The left tab is titled 'nginx-deployment.yaml' and contains the following YAML configuration:

```
! nginx-deployment.yaml ×
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: nginx-deployment
5   labels:
6     app: nginx
7 spec:
8   replicas: 2
9   selector:
10    matchLabels:
11      app: nginx
12   template:
13     metadata:
14       labels:
15         app: nginx
16     spec:
17       containers:
18         - name: nginx
19           image: nginx:1.25
20           ports:
21             - containerPort: 8080
```

The right tab is titled 'nginx-service.yaml' and contains the following YAML configuration:

```
! nginx-service.yaml ×
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: nginx-service
5 spec:
6   selector:
7     app: nginx
8   ports:
9     - protocol: TCP
10    port: 80
11    targetPort: 8080
12
```

minikube

## create minikube cluster

```
minikube start --vm-driver=hyperkit
```

```
kubectl get nodes
```

```
minikube status
```

```
kubectl version
```

## delete cluster and restart in debug mode

```
minikube delete
```

```
minikube start --vm-driver=hyperkit --v=7 --alsologtostderr
```

```
minikube status
```

## kubectl commands

```
kubectl get nodes
```

```
kubectl get pod
```

```
kubectl get services
```

```
kubectl create deployment nginx-depl --image=nginx
```

```
kubectl get deployment
```

```
kubectl get replicaset  
kubectl edit deployment nginx-depl
```

## **debugging**

```
kubectl logs {pod-name}  
kubectl exec -it {pod-name} -- bin/bash
```

## **create mongo deployment**

```
kubectl create deployment mongo-depl --image=mongo  
kubectl logs mongo-depl-{pod-name}  
kubectl describe pod mongo-depl-{pod-name}
```

## **delete deployment**

```
kubectl delete deployment mongo-depl  
kubectl delete deployment nginx-depl
```

## **create or edit config file**

```
vim nginx-deployment.yaml  
kubectl apply -f nginx-deployment.yaml  
kubectl get pod  
kubectl get deployment
```

## **delete with config**

```
kubectl delete -f nginx-deployment.yaml
```

#Metrics

kubectl top The kubectl top command returns current CPU and memory usage for a cluster's pods or nodes, or for a particular pod or node if specified.

Which kubectl apply : u can both create and update a component .

## YAML configuration file :

3-parts of k8s config-file .

## K8s YAML Configuration File - 1

- Also called "Kubernetes manifest"
- Declarative: A manifest **specifies the desired state** of a K8s component
- Config files are in **YAML format**, which is user-friendly, but **strict indentation!**
- Config files should be stored in version control

Each configuration file has **3 parts**:

demo-test-deployment.yaml 418 B

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: my-app
5   labels:
6     app: my-app
7 spec:
8   replicas: 2
9   selector:
10    matchLabels:
11      app: my-app
12   template:
13     metadata:
14       labels:
15         app: my-app
16     spec:
17       containers:
18         - name: my-app
19           image: my-image
20           env:
21             - name: SOME_ENV
22               value: $SOME_ENV
23           ports:
24             - containerPort: 8080
```

## 2) specification

- Attributes of "spec" are specific to the kind

The image shows two terminal windows side-by-side. The left window displays the contents of a file named 'nginx-deployment.yaml'. It includes lines 1 through 22. Lines 1, 2, and 3 are highlighted with a yellow box, showing 'apiVersion: apps/v1', 'kind: Deployment', and 'metadata:' respectively. Line 5, which starts 'spec:', is also highlighted with a yellow box. The right window displays the contents of a file named 'nginx-service.yaml'. It includes lines 1 through 12. Lines 1, 2, and 3 are highlighted with a yellow box, showing 'apiVersion: v1', 'kind: Service', and 'metadata:' respectively. Line 5, which starts 'spec:', is also highlighted with a yellow box.

```
! nginx-deployment.yaml x
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4 name: nginx-deployment
5 spec:
6 replicas: 2
7 selector: ...
8 template: ...
1 22

! nginx-service.yaml x
1 apiVersion: v1
2 kind: Service
3 metadata:
4 name: nginx-service
5 spec:
6 selector: ...
7 ports: ...
1 12
```

1. Metadata -> name of the component .
2. specification -> replica ,selector ,port
3. Status -> it will be automatically generated and added by k8s , the way it works is k8s will always compare what is the desired state and what is the actual state .

Where does k8s get that status data to automatically add in status part in yaml file >?  
Info comes from etcD - it hold the current status of k8s components .

**:: yaml validator to can check your yaml file here is it is correctly written as yaml is very strict in its indentations . And store these configuration file on with your code .**

### Blueprint for pods (template ) :

Deployments mange the pods that are below them

So how does it happen where does this whole thing defined in the configuration ?

Since it's an abstraction over pod , we have the pod configuration inside deployment configuration , template has its own metadata and spec section its a bulring for pod.

So this configuration apply to a pod means pod should have its own configuration .

### Connection components (labels and selectors and ports) :

In metadata part u give components of deployment a key- value pair .

Pod get the label through the template blueprint

This label is matched by a selector .

Selector will connect the server to the deployment or pod .

The diagram illustrates the relationship between deployment status and deployment configuration. On the left, a terminal window shows the deployment status with 1 replica. On the right, a terminal window shows the deployment configuration with 2 replicas. A red 'X' symbol is positioned between them, indicating a mismatch between the current status and the desired configuration.

3) status

- Automatically generated and added by Kubernetes
- K8s gets this information from etcd, which holds the current status of any K8s component

Desired status: 2 replica

```
! nginx-deployment.yaml ×  
1 apiVersion: apps/v1  
2 kind: Deployment  
3 metadata:  
4 | name: nginx-deployment  
5 | labels:--  
6 spec:  
7 | replicas: 2  
8 | selector:--  
9 | template:--  
10
```

```
% kubectl apply -f nginx-deployment.yaml
deployment.apps/nginx-deployment created
% kubectl get pod
NAME                               READY   STATUS    RESTARTS   AGE
nginx-deployment-8fd685b7d-8bjc2   1/1     Running   0          9s
nginx-deployment-8fd685b7d-gv7v2   1/1     Running   0          9s
% kubectl get service
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
kubernetes   ClusterIP  10.96.0.1   <none>        443/TCP   24h
nginx-service   ClusterIP  10.99.65.41  <none>        80/TCP    15s
```

```
% kubectl apply -f nginx-service.yaml
service/nginx-service created
% kubectl describe service nginx-service
Name:           nginx-service
Namespace:      default
Labels:          <none>
Annotations:    <none>
Selector:        app=nginx
Type:            ClusterIP
IP Family Policy: SingleStack
IP Families:    IPv4
IP:              10.99.65.41
IPs:             10.99.65.41
Port:            <unset>  80/TCP
TargetPort:      8080/TCP
Endpoints:      10.244.0.15:8080,10.244.0.16:8080
Session Affinity: None
Events:          <none>
```

Via above command u can see the updated status of the yaml file

```
% kubectl get pod -o wide
NAME          READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED LEADER
ESS GATES
nginx-deployment-8fd685b7d-8bjc2   1/1     Running   0          100s   10.244.0.15   minikube   <none>
nginx-deployment-8fd685b7d-gv7v2   1/1     Running   0          100s   10.244.0.16   minikube   <none>
% kubectl get deployment nginx-deployment -o yaml > nginx-deployment-result.yaml
```

## Project :

Web application requisition to its DB .

: first we will create a mongoDB pod

In order to talk to that pod we gonna need a service a internal service means no external requests from the internet are allowed to the pod .

Only component inside the same cluster can talk to it.

: second we will create a mongo Express deployment

One we need DB URL of mongo DB so that mongoExpress can connect to it .

Second one is credentials username and password so that to authentically to DB .

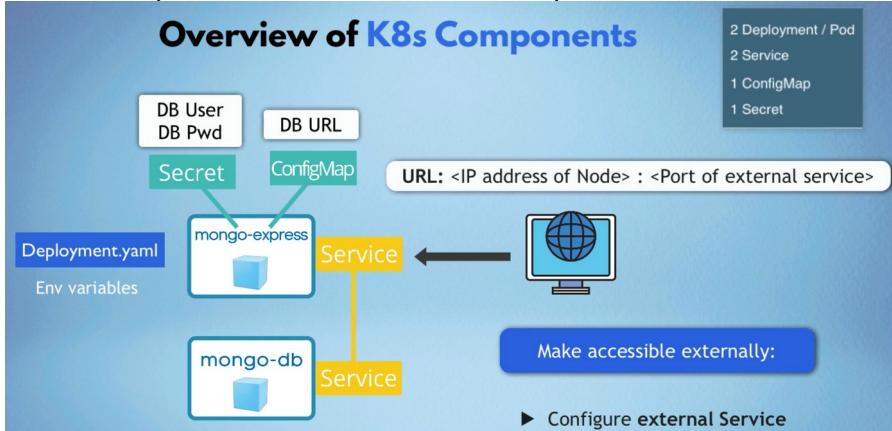
And we can pass these credential and DB url information to the mongoExpress through mongoExpress.yaml file(config\_file) through ENV\_variables.

Create a configMap that contains DB\_URL .

Create a secret that contains Credentials.

And we gonna reference both of these inside the config\_file .

So we need mongo express to be accessible through browser for that we need to create a external service. That would allow external request to talk to the pod . URL =><ip address of the node> : <port of external service>



### **Browser request flow through k8s components :**

Request come from the browser -> mongo express external service -> mongo\_express pod -> internal service of mongo\_DB(configMap= database URL) -> mongoDB pod and authenticate the request using the credentials .

#### 1. Create a mongo DB deployment config\_file

So to docker mongo images -> to see what are the external configuration or ports etc information about the image .

So this is a configuration file that will be checked in a repo -> so u don't want to show credential secrets in configuration file itself .

So we will create a secret from where we wanna take the values .

#### **Create secret**

**Generate -> base 64 for more security .**

: echo -n 'username' | base64

```
: echo -n 'password' | base64
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl apply -f mongo-secret.yaml
secret/mongodb-secret created
[(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl get secret
NAME          TYPE      DATA   AGE
mongodb-secret  Opaque    2      39s
```

Only when we apply secret then u can use that on your configuration file before applying configuring file .

Next is create in mongoDB internal service so that other components and pods can

## Service Configuration File

- ▶ **kind:** "Service"
- ▶ **metadata/name:** a random name
- ▶ **selector:** to connect to Pod through label
- ▶ **ports:**
  - port:** Service port
  - targetPort:** containerPort of Deployment

talk to that mongoDB.

```
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl apply -f mongo.yaml
deployment.apps/mongodb-deployment created
service/mongodb-service created
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl get service
NAME          TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
kubernetes   ClusterIP  10.96.0.1    <none>       443/TCP   153m
mongodb-service  ClusterIP  10.107.53.39  <none>       27017/TCP  60s
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl describe service mongodb-service
Name:            mongodb-service
Namespace:       default
Labels:          <none>
Annotations:    <none>
Selector:        app=mongodb
Type:           ClusterIP
IP Family Policy: SingleStack
IP Families:    IPv4
IP:              10.107.53.39
IPs:             10.107.53.39
Port:            <unset>  27017/TCP
TargetPort:      27017/TCP
Endpoints:      10.244.0.5:27017
Session Affinity: None
Events:          <none>
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl get pod -o wide
NAME                           READY   STATUS    RESTARTS   AGE     IP           NODE   NOMI
NATED NODE   READINESS GATES
mongodb-deployment-585bb4fddc-lrnrc  1/1     Running   0          3m17s  10.244.0.5   minikube  <non
e>          <none>
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubecrt get all | grep mongodb
zsh: command not found: kubecrt
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl get all | grep mongodb
pod/mongodb-deployment-585bb4fddc-lrnrc  1/1     Running   0          6m33s
service/mongodb-service  ClusterIP  10.107.53.39  <none>       27017/TCP  6m33s
deployment.apps/mongodb-deployment  1/1     1           1          6m33s
replicaset.apps/mongodb-deployment-585bb4fddc  1           1           1          6m33s
```

Next we will create a mongoExpress deployment and its external service. Also a external configuration where we will put Database url for mongoDB (configMap)

We need to tell it which database to connect to internal service mongoDB address .

Which credentials to authenticate >

# ConfigMap Configuration File

- ▶ **kind:** "ConfigMap"
- ▶ **metadata/name:** a random name
- ▶ **data:** the actual contents -  
in key-value pairs

Configmap must already be in k8s cluster when referencing it .

```
- name: ME_CONFIG_MONGODB_ADMINUSERNAME
  valueFrom:
    secretKeyRef:
      name: mongodb-secret
      key: mongo-root-username
- name: ME_CONFIG_MONGODB_ADMINPASSWORD
  valueFrom:
    secretKeyRef:
      name: mongodb-secret
      key: mongo-root-password
- name: ME_CONFIG_MONGODB_SERVER
  valueFrom:
    configMapKeyRef:
      name: mongodb-configmap
      key: database_url
```

So these are the three env variable that mongoExpress need to authenticate with mongoDB .

```
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl apply -f mongo-express.yaml
deployment.apps/mongo-express created
service/mongo-express-service created
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl apply -f mongo-configmap.yaml
configmap/mongodb-configmap created
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl get pod
NAME                               READY   STATUS    RESTARTS   AGE
mongo-express-6cfbc86cb6-jrd54     1/1     Running   0          66s
mongodb-deployment-585bb4fddc-lrnrc 1/1     Running   0          79m
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl logs mongo-express-6cfbc86cb6-jrd54
4
Mongo Express server listening at http://0.0.0.0:8081
Server is open to allow connections from anyone (0.0.0.0)
basicAuth credentials are "admin:pass", it is recommended you change this in your config.js!
```

Yes!! Mongo - Express server has started . And data base is connected .

Now final step is to access mongo express from the browser .  
So for that we need external service for mongo express .

Load balancer here accepts external requests assigning a service a external ip address

Node port -> port where this external ip address will be opened .  
Port that u need to put into the browsers to access this service

Finally two service we have created ->

## How to expose the Service externally?

► **type:** "LoadBalancer"

► **nodePort:** must be between 30000-32767

```
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % kubectl get service
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP  PORT(S)        AGE
kubernetes     ClusterIP  10.96.0.1    <none>       443/TCP       4h1m
mongo-express-service LoadBalancer  10.100.57.59  <pending>   8081:30000/TCP  10m
mongodb-service ClusterIP  10.107.53.39  <none>       27017/TCP     89m
```

Cluster ip will get service its internal ip address only reachable with in the cluster.  
Loadbalancer will also give internal ip address to the service but also give service a external ip address where the external request will be coming from .

In clusterip u have ports for only internal ip address where as we can see for loadbalancer service we have both 8081 port for internal and port 30000 for external ip address.

Currently we can see that for load balancer external service we have pending ip address ->

In minikube we have to do command : minikube service <serviceName> // this command will assign external service a public ip address.

```
(base) vishaldwivedi@vishals-MacBook-Air-2 1.K8sComponents % minikube service mongo-express-service
|-----|-----|-----|-----|
| NAMESPACE |     NAME     | TARGET PORT |     URL      |
|-----|-----|-----|-----|
| default  | mongo-express-service |      8081 | http://192.168.49.2:30000 |
|-----|-----|-----|-----|
🏃 Starting tunnel for service mongo-express-service.
|-----|-----|-----|-----|
| NAMESPACE |     NAME     | TARGET PORT |     URL      |
|-----|-----|-----|-----|
| default  | mongo-express-service |          | http://127.0.0.1:56528 |
|-----|-----|-----|-----|
⚠️ Opening service default/mongo-express-service in default browser...
! Because you are using a Docker driver on darwin, the terminal needs to be open to run it.
```

give a URL to external service in minikube -> minikube service mongo-express-service

